

Edeflip: Supervised Word Translation between English and Yoruba

Ikeoluwa Abioye

Ike.23@dartmouth.edu

Jiani Ge

Jiani.Ge.23@dartmouth.edu

Abstract

In recent years, embedding alignment has become the state-of-the-art machine translation approach. It can yield high translation precision without parallel corpora. However, extant research and application of embedding alignment mostly focus on high-resource languages with high-quality monolingual embeddings. It is unclear if and how low-resource languages may similarly benefit. In this study, we implement an established supervised embedding alignment method on English and Yoruba, the latter a low-resource language. We found that higher embedding quality (size, correct diacritical notations) and applying embedding normalization procedures increase word translation accuracy, with, additionally, an interaction effect between the two. Our results demonstrate the limit of supervised embedding alignment when it comes to low-resource languages, for which there are additional factors that need to be taken into consideration, such as the importance of curating high-quality monolingual embeddings, in order to produce high-quality translations. We hope our work will be a starting point for further research on the embedding alignment that takes into account the challenges that low-resource languages face.

1 Introduction

In recent years, alignment of monolingual embedding spaces has replaced training on parallel corpora as the state-of-the-art machine translation approach (Jansen, 2018; Conneau et al., 2017). Embeddings are representations of language tokens (words, sub-words, phrases) in a high-dimensional vector space. They are generated through an algorithm that learns the context in which the tokens in the training material appear, thereby representing the semantic and syntactic relationships between the tokens. As similar semantic and syntactic relationships exist in all languages, embedding alignment exploits this regularity and learns

a transformation matrix that matches the relative positions of tokens with similar meanings in their respective monolingual embedding spaces (Jansen, 2018; Conneau et al., 2017). As of today, continued refinement of the alignment algorithm has led to translation performance that either matches or exceeds that of the previous models based on parallel corpora (Conneau et al., 2017; Artetxe et al., 2019).

Embedding alignment has the advantage of reducing the need for multilingual parallel corpora, which are costly to curate and limited in availability for many language pairs. It also makes multilingual (rather than bilingual) translation more practical, as the embeddings of multiple languages can be mapped onto the same vector space using only monolingual resources. However, most of the research and published resources for embedding alignment to-date are for high-resource languages, such as English, Spanish, or Chinese. Even though the reduced need for parallel corpora promises increased accessibility for the translation of low-resource languages, it is unclear how they would actually fare under this approach.

The present study addresses this question by focusing on word translation between English and Yoruba. Yoruba is a language of the Niger-Congo family with over 50 million speakers in West Africa, though under-resourced in NLP research (Joshi et al., 2020). As Alabi et al. (2020) points out, the available monolingual resources for Yoruba, such as the pre-trained fastText embeddings, are of very low quality. First, many tokens in these embeddings are missing diacritics or marked incorrectly. In Yoruba, diacritics mark distinct letters and tones and are crucial to meaning representation. Second, these embeddings contain many English or other non-Yoruba tokens, which may bias the vector space. Finally, they are many times smaller in size (the pre-trained Wikipedia

fastText Yoruba embedding contains 21,730 tokens, whereas the English one contains more than 2 million (Mikolov et al., 2018)).

As embedding alignment relies on monolingual embeddings, it is unclear how small and noisy embeddings, which are often the case for low-resource languages, affect translation performance. In this study, we adopt the supervised embedding alignment method from (Conneau et al., 2017), manipulate the quality of the Yoruba embedding, and examine its effect on word translation precision. Investigating how the available embeddings for low-resource languages perform in embedding alignment may help us assess the extent to which the current state-of-the-art machine translation approach immediately benefits these languages, as well as the factors that need to be addressed to obtain high-quality machine translation results.

2 Methodology

2.1 Data

We use data for three purposes in the study. All data are collected from online, published sources.

Monolingual embeddings: We use existing English and Yoruba embeddings for the alignment. We use a pretrained fastText embedding for English, with 1 million word vectors trained on Wikipedia and web crawl data (Mikolov et al., 2017). For Yoruba, we use a pretrained fastText embedding, with 21,730 word vectors trained on Wikipedia (Bojanowski et al., 2017), and a second embedding trained by Alabi et al. (2020), with 113,572 word vectors trained on high-quality texts from Yoruba news and websites with complete diacritics notations. We use word embeddings, rather than embeddings with subword tokens, because the latter is not available for the high-quality Yoruba corpora.

Ground-truth dictionary: For the ground-truth dictionary used in supervised alignment, we scrape the online English-Yoruba dictionary Alabi et al. (2020) and produce 3,693 unique word pairs. We create separate entries for words with more than one translation. Entries with multi-word translations are excluded, as the embeddings we use only contain one-word tokens.

Evaluation dictionary: For evaluating translation precision, we take the word pairs from the WordSim353 dataset translated into Yoruba by (Al-

abi et al., 2020). We remove entries with multi-word translations and out-of-vocabulary words, and obtain 161 unique word pairs.

2.2 Algorithm

We adapt our algorithm from Facebook MUSE research project (Conneau et al., 2017) on supervised machine translation. While Yoruba is not included in the original research, we decided to adapt the algorithm for our task and see how well it performed. Conneau et al. (2017) most recently developed an unsupervised alignment method that uses adversarial learning to create a synthetic parallel dictionary, then applying supervised Procrustes alignment to the embeddings, using the synthetic dictionary as anchor points. Procrustes alignment matches the relative positions of two embedding spaces by finding the optimal orthogonal matrix that rotates, scales, and transforms them to be mostly closely aligned. The Procrustes problem has a closed-form solution through Singular Value Decomposition (SVD):

$$W^* = \arg \max_{W \in O_d(R)} \|WX - Y\|_F = UV^T$$

Where $U\Sigma V^T = \text{SVD}(YX^T)$, W is the mapping or transformation matrix, X is the source language embedding space, and Y is the target language embedding space.

Due to the high computational demand of the adversarial learning phase of the method, we decide to perform Procrustes alignment using a human-created ground-truth dictionary as anchor points, thus rendering the overall process supervised. In half of the trials, prior to the alignment, we center and normalize the embeddings using the tensor operations from the PyTorch package, in order to make the distributions of the two embedding spaces more similar.

2.3 Experiment

We create four pairs of aligned English and Yoruba embeddings by manipulating two variables. Two of the alignments normalize the embeddings before alignment while the other two do not; and two of the alignments use the Wikipedia Yoruba embedding, while the other two the curated Yoruba embedding. Hereafter, the four conditions will be referred to as **wiki-unnormalized**, **wiki-normalized**, **curated-unnormalized**, and **curated-normalized**. We evaluate the alignment result by calculating the English to Yoruba word translation precision based

Precision@	k = 1	k = 5	k = 10
Wiki-unnorm	10.56	18.01	21.74
Wiki-norm	12.42	19.25	25.47
Curated-unnorm	6.88	11.88	14.38
Curated-norm	19.38	23.75	29.38

Table 1: en-yo word translation precision at k=1, 5, 10

on the evaluation dataset. As word translation is generated by retrieving the nearest neighbors of the source word in the target language embedding space, we calculate the probabilities that the top 1, 5, and 10 retrievals (k=1, 5, and 10) contain the correct translation. For example, a precision of 0.3 at k=5 means that for 30 percent of source words, the correct translation (or one of the correct translations) appears in the top five nearest neighbors in the target language embedding space.

3 Results

Table 1 shows the word translation precision of the four alignment conditions. Overall, the wiki-unnormalized condition consistently yields the lowest precision, whereas the curated-normalized condition the highest. Overall, the normalized alignments yield a higher precision than the unnormalized alignments that use the same embeddings. Rather than a main effect of embedding quality, however, our results show an interaction between embedding quality and normalization: whereas embedding normalization greatly improved the word translation precision for the curated conditions, the improvement for the wiki conditions is relatively small.

We also assessed the word translation results qualitatively by looking at individual queries and visualizing the aligned embeddings using t-SNE plots (Figure 1). The outcomes are consistent with the above-reported normalization and interaction effects. For example, given the source word ‘sea’, only the curated-normalized alignment yields the correct Yoruba translation ‘òkun’ (at k=3). The top retrievals of the wiki alignments are similar to each other and include omí ‘water’, irinàjò ‘travel’, which are arguably related to ‘sea’. The top retrievals of the curated-unnormalized model, however, bares no visible semblance to the meaning of ‘sea’; they include, tápárebì ‘old age’ and kápòṇónlẹ ‘the only one.’ As can be seen in Figures 1-4, before normalization, the English and

Yoruba embeddings are on different scales, with the Yoruba vectors being much farther apart. After normalization, the vectors of the two languages are on a similar scale, though they still do not constitute a perfect mapping.

4 Discussion

4.1 Word translation performance with en-yo

Overall, the results are as expected. A comparison of the four alignment conditions shows clear effects of both the quality of the embeddings and the embedding normalization procedure. Using larger and correctly marked Yoruba embeddings and normalizing the English and Yoruba embeddings before alignment increases the quality of the alignment and hence word translation accuracy. This is consistent with the fact that the precisions we have obtained are significantly lower than those in the published articles. For example, in [Conneau et al. \(2017\)](#), the same supervised alignment method produces a word translation accuracy at k=1 of 77.4 for English-Spanish, 68.4 for English-German, and 40.6 for English-Chinese (p. 6). For one thing, these are high-resource languages, and the embeddings that these alignments use are multitudes larger than the Yoruba embeddings that we use. Further, the low precision that we obtain for English-Yoruba may be in part a extension of the trend that precisions are lower for language pairs that are more morpho-syntactically different, which are also supported by the just-mentioned results for Spanish, German, and Chinese.

The interaction effect between embedding quality and normalization was not predicted, though in hindsight comprehensible. We reason that it is because the English and Yoruba embeddings trained on Wikipedia have a more similar shape from the start, as they are both trained using fastText library with similar parameters, and their training corpora (Wikipedia, among others) are likely to have similar words and word frequencies. In comparison, the curated Yoruba embedding is trained independently on more distinct ([Alabi et al., 2020](#)). A similar phenomenon is reported in [Conneau et al. \(2017\)](#), where using the same alignment method, the word translation accuracy between English and Italian is higher when Wikipedia embeddings are used for both languages, compared to using Wikipedia embeddings for one language and CBOW embeddings for another (p. 8). Nevertheless, in the curated-normalized condition in our study, the possible

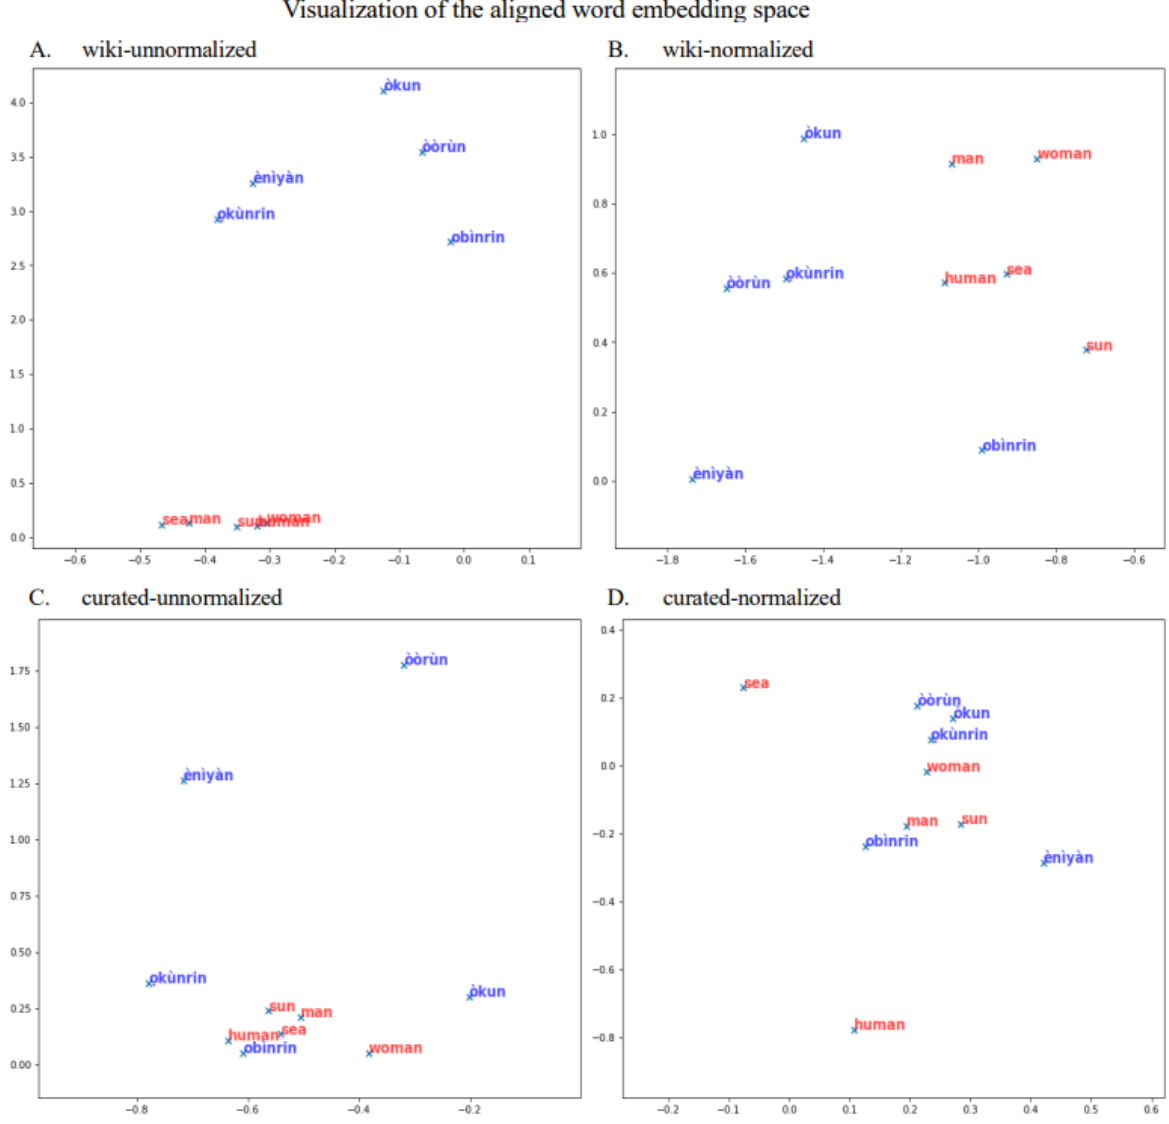


Figure 1: t-SNE plots for the aligned en-yo embeddings.

“disadvantage” of dissimilarity between source and target embeddings is offset by the normalization procedure and the superior quality of the curated Yoruba embeddings.

Taken together, our study shows that for word translation through embedding alignment, the quality (size and correctness) of the embeddings, normalization procedure, and the similarity between the original source and target embeddings all contribute to translation precision. While these factors may be less salient for translations between languages that have comparable high-quality embeddings (e.g., English and Spanish), they constitute crucial challenges that need to be overcome for projects that involve low-resource languages.

4.2 Ethical considerations

Our study also brings to attention several ethical considerations, primarily due to Yoruba’s status as a low-resource language. Limited Yoruba NLP resource can inadvertently lead to exclusion and demographic bias. Compromised machine translation programs performance may risk poor translations being circulated without adequate verification. Such inaccuracies could result in miscommunication or misinterpretation, potentially leading to conflict and harm, especially considering the fact that while a translation may be technically accurate, it could also be culturally inappropriate or offensive. Furthermore, a significant disparity exists in the availability of computational resources; large organizations with substantial processing power

can achieve superior results compared to smaller groups with limited resources. There is also an increased risk of privacy violations and exploitation of a community’s linguistic resources without informed consent or tangible benefits.

5 Conclusion

Our study has shown that monolingual word embedding quality and normalization procedure affects the word translation performance of embedding alignment. Our findings identify challenges for machine translation with low-resource languages under the current state-of-the-art approach. Admittedly, the present study has many limitations. Future work may enhance the translation performance by using a larger training dictionary, incorporating more diverse tasks and metrics for evaluation, and, ideally, working with embeddings with even higher quality. We hope our work may be the starting point for a machine translation research that is more sensitive to the factors and challenges that particularly impact low-resource languages, and for efforts to address the resource inequality in the NLP.

References

- Jesujoba O. Alabi, Kwabena Amponsah-Kaakyire, David I. Adelani, and Cristina España-Bonet. 2020. [Massive vs. curated word embeddings for low-resourced languages. the case of yorùbá and twi.](#)
- Mikel Artetxe, Gorka Labaka, and Eneko Agirre. 2019. [Bilingual lexicon induction through unsupervised machine translation.](#) In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 5002–5007, Florence, Italy. Association for Computational Linguistics.
- Piotr Bojanowski, Edouard Grave, Armand Joulin, and Tomas Mikolov. 2017. [Enriching word vectors with subword information.](#)
- Alexis Conneau, Guillaume Lample, Marc’Aurelio Ranzato, Ludovic Denoyer, and Hervé Jégou. 2017. Word translation without parallel data. *arXiv preprint arXiv:1710.04087*.
- Stefan Jansen. 2018. [Word and phrase translation with word2vec.](#)
- Pratik Joshi, Sebastin Santy, Amar Budhiraja, Kalika Bali, and Monojit Choudhury. 2020. [The state and fate of linguistic diversity and inclusion in the NLP world.](#) In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 6282–6293, Online. Association for Computational Linguistics.
- Tomas Mikolov, Edouard Grave, Piotr Bojanowski, Christian Puhersch, and Armand Joulin. 2017. [Advances in pre-training distributed word representations.](#)
- Tomas Mikolov, Edouard Grave, Piotr Bojanowski, Christian Puhersch, and Armand Joulin. 2018. [Advances in pre-training distributed word representations.](#) In *Proceedings of the International Conference on Language Resources and Evaluation (LREC 2018)*.